


CSEG601 & CSE5601

Spatial Data Management & Application :



Spatial Joins


Sungwon Jung

Big Data Processing and Database Lab.
Dept. of Computer Science and Engineering
Sogang University
Seoul, Korea
Tel: +82-2-705-8930
Email : jungsung@sogang.ac.kr

1




R-tree based Join



2

Join1(R,S)

- Tree synchronized traversal algorithm
Join1(R,S)
Repeat
 Find a pair of intersecting entries E in R and F in S
 If R and S are leaf pages then
 add (E,F) to result-set
 Else Join1(E,F)
Until all pairs are examined




3

CPU – Time Tuning

- Two ways to improve CPU – time
 - Restricting the search space
 - Spatial sorting and plane sweep

4

Reducing CPU bottleneck



5

Join2(R,S,IntersectedVol)

Join2(R,S,IV)

Repeat

Find a pair of intersecting entries E in R and F in S that overlap with IV

If R and S are leaf pages then

add (E,F) to result-set

Else Join2(E,F,CommonEF)

Until all pairs are examined


- In general, number of comparisons equals
 - $\text{size}(R) + \text{size}(S) + \text{relevant}(R) * \text{relevant}(S)$
- Reduce the product term

6

Restricting the search space

Join1: 7 of R * 7 of S

= 49 comparisons





Now: 3 of R * 2 of S
= 6 comp

Plus Scanning:
7 of R + 7 of S
= 14 comp

7

Using Plane Sweep

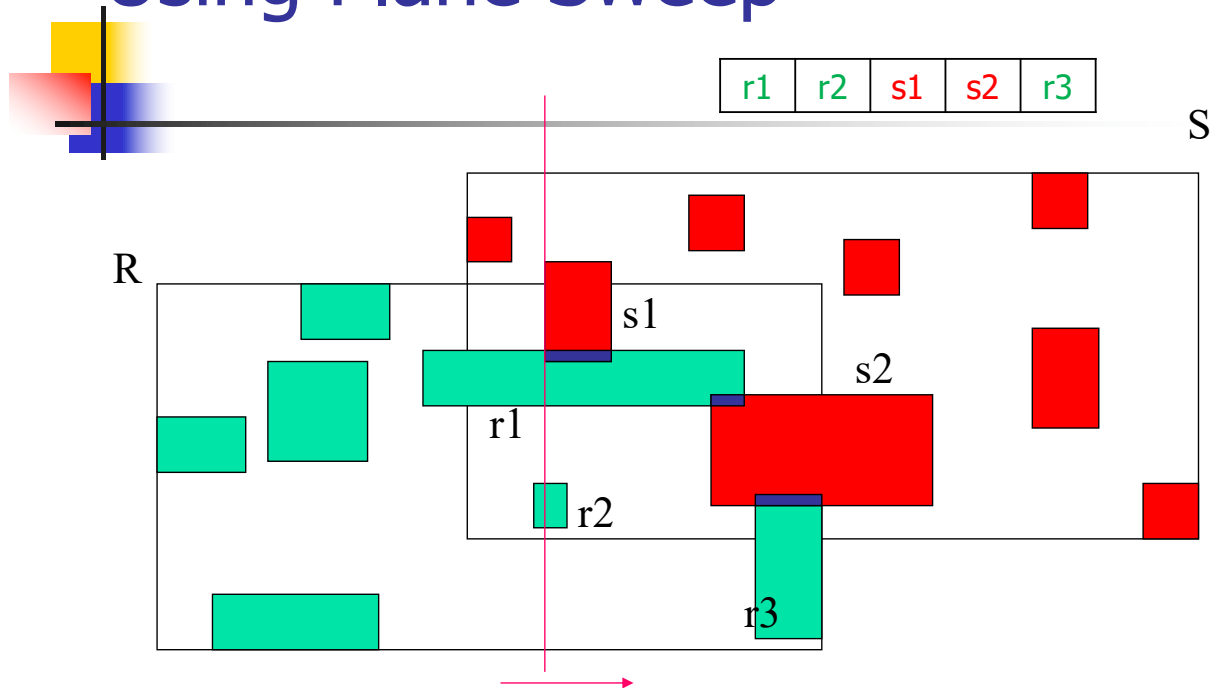
Sort the set of rectangles w.r.t. the x values of their lower-left corners in the intersected area of R and S



Consider the extents along x-axis
Start with the first entry r1
sweep a vertical line

8


Using Plane Sweep



Check if $(r1, s1)$ intersect along y-dimension
Add $(r1, s1)$ to result set

9

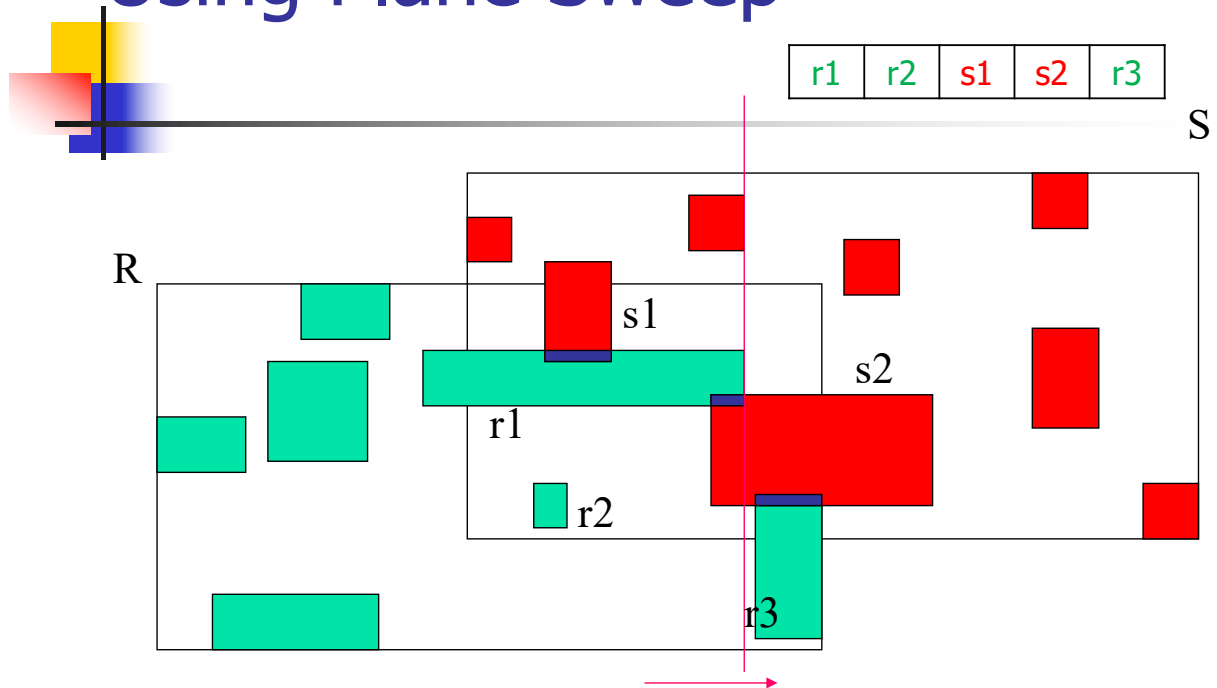
Using Plane Sweep



Check if $(r1, s2)$ intersect along y-dimension
Add $(r1, s2)$ to result set

10


Using Plane Sweep



Reached the end of $r1$
Start with next entry $r2$

11

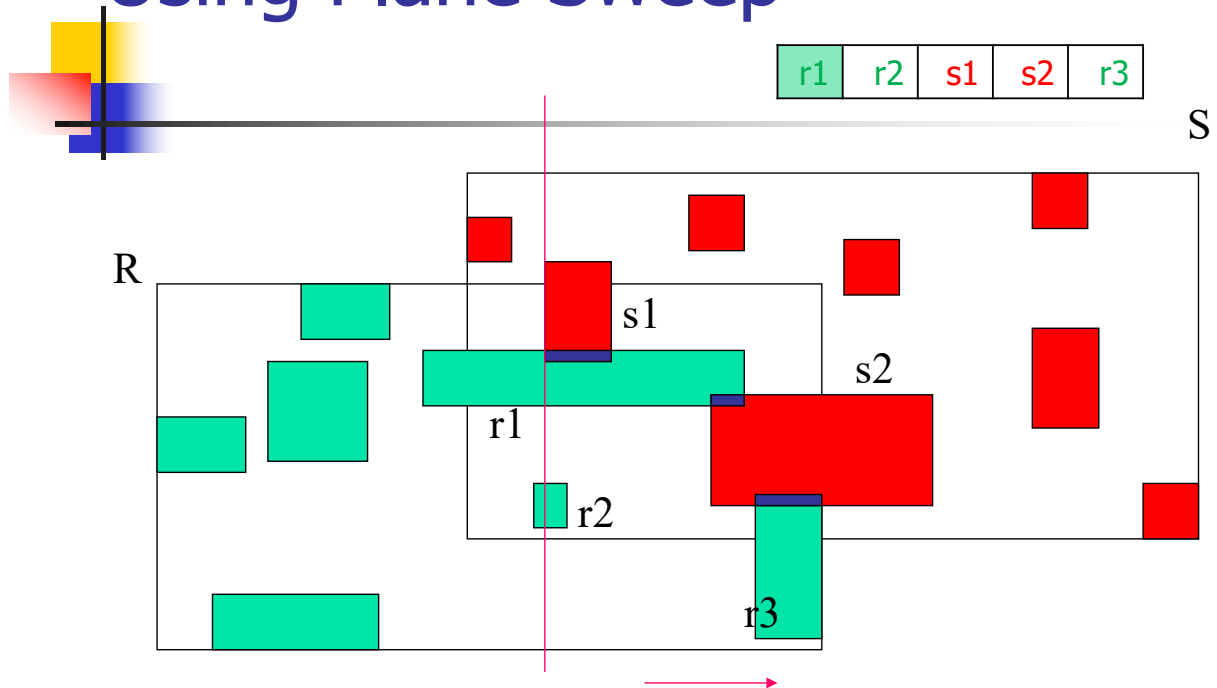
Using Plane Sweep



Reposition sweep line

12

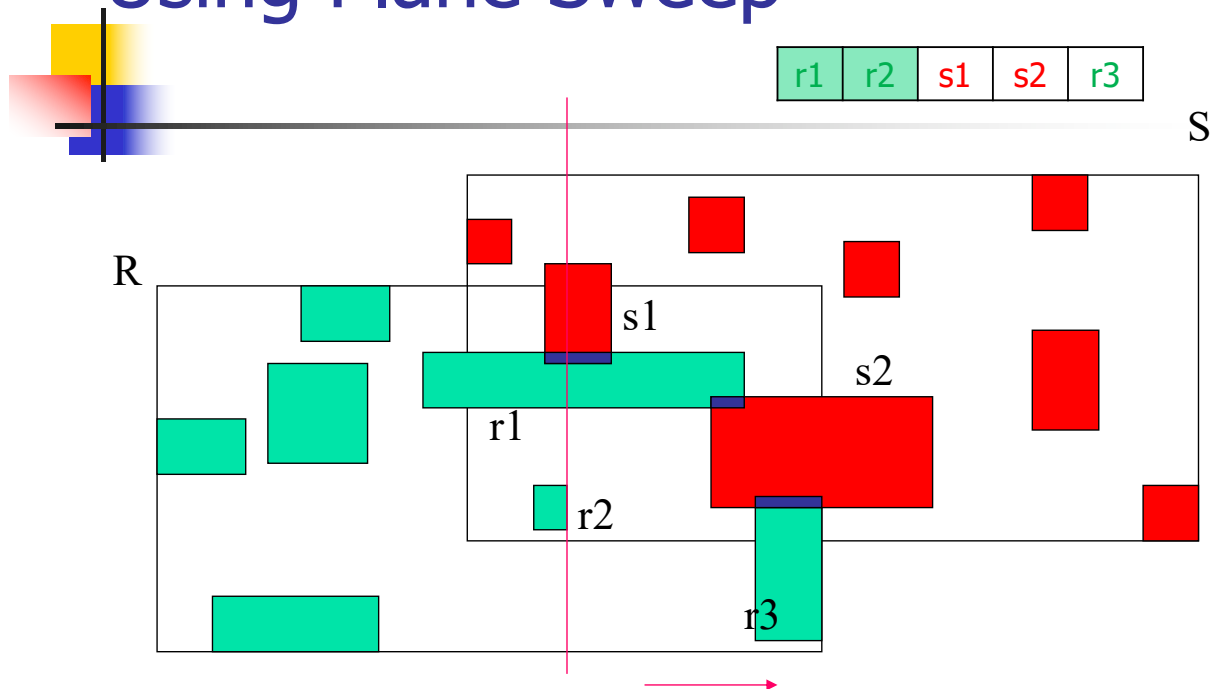
Using Plane Sweep



Check if r2 and s1 intersect along y
Do not add (r2,s1) to result

13


Using Plane Sweep



Reached the end of r2
Start with next entry s1

14

Using Plane Sweep



Total of $2(r1) + 1(r2) + 0(s1) + 1(s2) + 0(r3) = 4$ comparisons

15

Spatial Join Algorithm using two R trees with different depths


```

SJ(R1, R2:R_NODE);
01  BEGIN
02      FOR (all Er2 in R2) DO
03          FOR (all Er1 in R1) DO
04              IF (overlap(Er1.rect, Er2.rect)) THEN
05                  IF (R1 and R2 are leaf pages) THEN
06                      output(Er1.oid, Er2.oid)
07                  ELSE IF (R1 is a leaf page) THEN
08                      ReadPage(Er2.ptr);
09                      SJ(Er1.ptr, Er2.ptr)
10                  ELSE IF (R2 is a leaf page) THEN
11                      ReadPage(Er1.ptr);
12                      SJ(Er1.ptr, Er2.ptr)
13                  ELSE
14                      ReadPage(Er1.ptr), ReadPage(Er2.ptr);
15                      SJ(Er1.ptr, Er2.ptr)
16                  END-IF
17              END-IF
18          END-FOR
19      END-FOR;
20  END.
    
```


16

Spatial Hash Join

- Hash join based on Space-driven structures (with redundancy)




- Hash join based on Data-driven structures (with overlapping)



17

Spatial Hash Join based on overlapping

- Step 1: **Initial partitioning**
 - Partition R
 - Each bucket has roughly same number of rectangles
 - Each bucket should fit in memory
 - The overlapping of bucket extents is minimized




- Content of the buckets

$$A = \{1, 14, 10\}, \quad B = \{2, 4, 7, 11\},$$

$$C = \{3, 8, 9\}, \quad D = \{5, 6, 12, 13\}$$


18



Spatial Hash Join based on overlapping

■ Step 2: **Second partitioning**

- Partition S
- Keep same buckets from step 1
- Assign rectangle of S to any bucket whose extent overlaps it



- Content of the buckets

$$\begin{aligned} A' &= \{c, d\}, & B' &= \{a, b, d, g\}, \\ C' &= \{e, g, h\}, & D' &= \{a, e\} \end{aligned}$$


19




Spatial Hash Join based on overlapping

■ Step 3: **Join phase**

- From step 1 and 2, we have two sets of buckets
- Each bucket from R and S has same extent and location
- Test each rectangle from bucket where *plane-sweeping* algorithm can be used




(a)



(b)


- [A, A'], [B, B'], [C, C'], [D, D'] must be joined

20



Spatial Hash Join based on redundancy

- Partition based Spatial Merge Join
 - Both sets R and S are partitioned with replication
 - Space is regularly tiled
 - Partitions either correspond to tiles or are determined from them using hashing

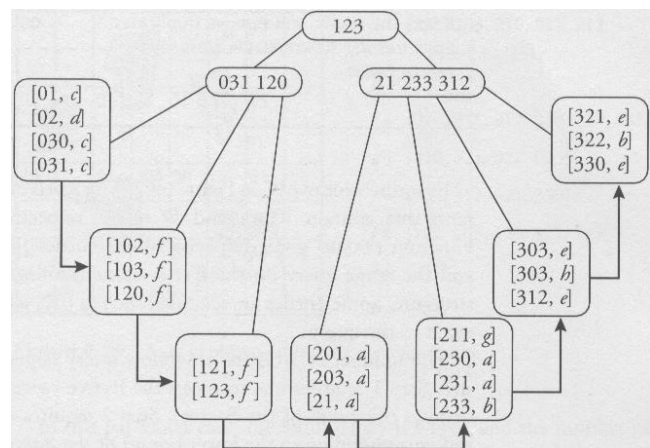



21




z-Ordering Spatial Join

- z-ordering tree



22



z-Ordering Spatial Join


- Assumptions
 - The leaves of each z-ordering tree represent a list L of entries of the form $[z, oid]$, ordered on z
 - oid : id of an object whose approximation contains the cell with z -order value (key) z
 - One cell with key z is contained in a cell with key z' if z' is a prefix of z
- Basic Concept of Join algorithm
 - The lists of entries corresponding to the two relations, $L1$ and $L2$, are merged.
 - Keep one pair of entries from the two lists as a candidate for the refinement step if one key is a prefix of the other.
 - Candidate pairs have to be sorted in order to remove duplicates.

23




z-Ordering Spatial Join

- Notations
 - Rectangles **above (below)** the z axis represent cells of list $L1(L2)$.
 - **Lowest (Largest)** value of a rectangle on the z axis is the key z (z') associated with the cell c .
 - $c : z' = scc(z)$ (smallest cell in the lower-right corner of z)
 - If $z = 30$, (a maximal depth of 7 for the decomposition) $z' = 3033333$.





z-Ordering Spatial Join



ZORDERINGJOIN (L_1, L_2 : list of *ids*): set of pairs of entries

begin

result: Set of pairs of ids, initially empty

while not (eof(L_1) **and** empty(S_1) **and** eof(L_2) **and** empty(S_2))

begin

event = MIN (CURRENT(L_1), SCC(top(S_1)),

CURRENT(L_2), SCC(top(S_2)))

if (*event* = CURRENT(L_1)) **then** // left bound of a rectangle

ENTRY (L_1 , S_1)

else if (*event* = SCC(top(S_1))) **then** // right bound of a rectangle

result += EXIT (S_1 , S_2)

else if (*event* = CURRENT(L_2)) **then** // left bound of a rectangle

ENTRY (L_2 , S_2)

else if (*event* = SCC (top(S_2))) **then** // right bound of a rectangle

result += EXIT (S_2 , S_1);

end while

sort *result*; remove duplicates;

return result

end

	C_1	S_1	C_2	S_2	Event Actions
Step 0	A	()	D	()	
Step 1	B	(A)	D	()	event = current (L_1) = A
Step 2	B	(A)	E	(D)	event = current (L_2) = D
Step 3	B	(A)	E	()	event = scc(top(S_2)) = D result={[A,D]}
Step 4	C	(B,A)	E	()	event = current (L_1) = B
step 5	C	(B,A)	F	(E)	event = current (L_2) = E
Step 6	C	(B,A)	G	(F,E)	event = current (L_2) = F
Step 7	C	(A)	G	(F,E)	event = scc(top(S_1)) = B result= {[A,D]}, [B,F], [B,E]}
Step 8	C	()	G	(F,E)	event = scc(top(S_1)) = A result= {[A,D]}, [B,F], [B,E]} + {[A,F], [A,E]}